

LCIM: Mining Low Cost High Utility Itemsets

M. Saqib Nawaz, Philippe Fournier-Viger,
Naji Alhusaini, Yulin He, Youxi Wu, Debdatta
Bhattacharya

Introduction

- Huge amount of data
- A need to **understand** the data.
- **Pattern mining:** using algorithms to identify interesting patterns in data that can be also interpretable.
- **Frequent itemset mining:** finding patterns that appear frequently in the data.
- Recently, much attention on **high utility itemset mining** (HUIM), where the goal is to find values that co-occur in data and have a high utility (importance).

High Utility Itemset Mining

Input: Transaction database + Utility table + Minimum utility threshold (*minutil*)

Transaction database D

TID	Items
T_1	(a,1), (c,1)
T_2	(e,1)
T_3	(a,1), (b,5), (c,1), (d,3), (e,1)
T_4	(b,4), (c,3), (d,3), (e,1)
T_5	(a,1), (c,1), (d,1)
T_6	(a,2), (c,6), (e,2)
T_7	(b,2), (c,2), (e,1)

Utility table

Item	Unit profit
a	5\$
b	2\$
c	1\$
d	2\$
e	3\$

minutil threshold

minutil = 30\$

Output: High Utility Itemsets (HUIs),
i.e., itemsets having utility $\geq \text{minutil}$

{b,d}:30\$	{a,c}:34\$
{b,e}:31\$	{b,d,e}:36\$
{a,c,e}:31\$	{b,c,e}:37\$
{b,c,d}:34\$	{b,c,d,e}:40\$

Limitations

- The focus in HUIM is on the **utility** of patterns (benefits).
- However, the cost or effort required to obtain these benefits is ignored.
- May find patterns that have:
a high utility but a very high cost
- May miss patterns that have:
a low cost but a relatively high utility

The proposal

- A novel problem named **Low-Cost High Utility Mining** that integrates the concept of **utility** with that of **cost**.
- **Goal:** Find itemsets that provide utility at a low cost.
- **Cost:** money, time resources consumed, effort

Example applications

- **E-learning:**
 - **Transaction:** list of activities taken by a student
 - **Cost:** the study time for each learning activity
 - **Utility:** the grades obtained at the final exam
 - **Low-cost high utility itemsets:** the learning activities that requires a small effort to obtain high grades
- **Medical pathway data:**
 - **Transaction:** list of treatments taken by a patient
 - **Cost:** the time or money spent for each treatment
 - **Utility:** the result of finally being cured or not
 - **Low-cost high utility itemsets:** the treatments that have a small cost and may lead to be cured.

How to combine utility and cost?

- **Simple solution:** subtract the cost from the utility and then apply HUIM algorithms.
 - **Problem:** utility and cost may be expressed in different units (e.g. hours, grades).
- **Solution in this paper:**
 - The cost and utility are modeled separately
 - Allows to use different units
 - Allows to calculate the average utility and average cost, which are more meaningful than total utility and total cost
(e.g. average time spent by student is 30 min)

Need to design a new algorithm!

Related work

- Algorithms for frequent itemset mining:
[Apriori](#), [FPGrowth](#), [Eclat](#), etc.
- Algorithms for high utility itemset mining
[UP-Growth](#), [UI-Miner](#), [FHM](#), [EFIM](#), [REX](#), etc.
 - No algorithm has integrated the concept of **cost**.
- Algorithms for finding low-cost sequential patterns:
[CorCEPB](#), [CEPN](#), [CEPB](#).
 - Analyzing **sequences** is different from analyzing transactions to find itemsets.
 - New algorithms must be designed.

Problem definition

Input: A transaction database with utility and cost values

Transaction	Items (<i>item, cost</i>)	Utility
T_1	($a, 5$), ($b, 10$), ($c, 1$), ($d, 6$), ($e, 3$), ($f, 5$)	40
T_2	($b, 8$), ($c, 3$), ($d, 6$), ($e, 3$)	20
T_3	($a, 5$), ($c, 1$), ($d, 2$)	18
T_4	($a, 10$), ($c, 6$), ($e, 6$), ($g, 5$)	37
T_5	($b, 4$), ($c, 2$), ($e, 3$), ($g, 2$)	21

Problem definition

Input: A transaction database with utility and cost values

Transaction	Items (<i>item, cost</i>)	Utility
T_1	$(a, 5), (b, 10), (c, 1), (d, 6), (e, 3), (f, 5)$	40
T_2	$(b, 8), (c, 3), (d, 6), (e, 3)$	20
T_3	$(a, 5), (c, 1), (d, 2)$	18
T_4	$(a, 10), (c, 6), (e, 6), (g, 5)$	37
T_5	$(b, 4), (c, 2), (e, 3), (g, 2)$	21

Let there be a set of items $I = \{i_1, i_2, \dots, i_n\}$

(e.g. learning activities a, b, c, d, e, f, g)

An itemset X is a set of items. e.g. {a,b}

Problem definition

Input: A transaction database with utility and cost values

Transaction	Items (<i>item, cost</i>)	Utility
T_1	$(a, 5), (b, 10), (c, 1), (d, 6), (e, 3), (f, 5)$	40
T_2	$(b, 8), (c, 3), (d, 6), (e, 3)$	20
T_3	$(a, 5), (c, 1), (d, 2)$	18
T_4	$(a, 10), (c, 6), (e, 6), (g, 5)$	37
T_5	$(b, 4), (c, 2), (e, 3), (g, 2)$	21

A **database** is a set of transactions $D = \{T_1, T_2, \dots, T_m\}$

Problem definition

Input: A transaction database with utility and cost values

Transaction	Items (<i>item, cost</i>)	Utility
T_1	($a, 5$), ($b, 10$), ($c, 1$), ($d, 6$), ($e, 3$), ($f, 5$)	40
T_2	($b, 8$), ($c, 3$), ($d, 6$), ($e, 3$)	20
T_3	($a, 5$), ($c, 1$), ($d, 2$)	18
T_4	($a, 10$), ($c, 6$), ($e, 6$), ($g, 5$)	37
T_5	($b, 4$), ($c, 2$), ($e, 3$), ($g, 2$)	21

e.g. Student #5 did learning activity « **b** » , « **c** » , « **e** » , « **g** »

Each **transaction** T_i is a set of items ($T_i \subseteq I$)

Problem definition

Input: A transaction database with utility and cost values

Transaction	Items ($item, cost$)	Utility
T_1	$(a, 5), (b, 10), (c, 1), (d, 6), (e, 3), (f, 5)$	40
T_2	$(b, 8), (c, 3), (d, 6), (e, 3)$	20
T_3	$(a, 5), (c, 1), (d, 2)$	18
T_4	$(a, 10), (c, 6), (e, 6), (g, 5)$	37
T_5	$(b, 4), (c, 2), (e, 3), (g, 2)$	21


e.g. student #5 took 4 minutes to do activity « b ».

Each **item** i in a **transaction** T has a **cost value** $c(i, T)$ that is a positive number

Problem definition

Input: A transaction database with utility and cost values

Transaction	Items (<i>item, cost</i>)	Utility
T_1	($a, 5$), ($b, 10$), ($c, 1$), ($d, 6$), ($e, 3$), ($f, 5$)	40
T_2	($b, 8$), ($c, 3$), ($d, 6$), ($e, 3$)	20
T_3	($a, 5$), ($c, 1$), ($d, 2$)	18
T_4	($a, 10$), ($c, 6$), ($e, 6$), ($g, 5$)	37
T_5	($b, 4$), ($c, 2$), ($e, 3$), ($g, 2$)	21



e.g. student #5 has received 21 points at the final exam

Each **transaction T** has a **utility value $u(T)$** that is a positive number representing the benefits

Problem definition

Input: A transaction database with utility and cost values

Transaction	Items ($item, cost$)	Utility
T_1	$(a, 5), (b, 10), (c, 1), (d, 6), (e, 3), (f, 5)$	40
T_2	$(b, 8), (c, 3), (d, 6), (e, 3)$	20
T_3	$(a, 5), (c, 1), (d, 2)$	18
T_4	$(a, 10), (c, 6), (e, 6), (g, 5)$	37
T_5	$(b, 4), (c, 2), (e, 3), (g, 2)$	21

Three parameters: $\text{minsup} > 0$, $\text{minutil} > 0$, $\text{maxcost} > 0$

Output: all the **low-cost high utility itemsets**

Each itemset X such that

its average utility $\text{au}(X) \geq \text{minutil}$,

its average cost $\text{ac}(X) \leq \text{maxcost}$,

its support $s(X) \geq \text{minsup}$.

Problem definition

Input: A transaction database with utility and cost values

Transaction	Items (<i>item, cost</i>)	Utility
T_1	(a, 5), (b, 10), (c, 1), (d, 6), (e, 3), (f, 5)	40
T_2	(b, 8), (c, 3) (d, 6), (e, 3)	20
T_3	(a, 5), (c, 1), (d, 2)	18
T_4	(a, 10), (c, 6), (e, 6), (g, 5)	37
T_5	(b, 4), (c, 2) (e, 3), (g, 2)	21

Example: $X = \{b,c\}$

Problem definition

Input: A transaction database with utility and cost values

Transaction	Items (<i>item, cost</i>)	Utility
T_1	$(a, 5), (b, 10), (c, 1), (d, 6), (e, 3), (f, 5)$	40
T_2	$(b, 8), (c, 3) (d, 6), (e, 3)$	20
T_3	$(a, 5), (c, 1), (d, 2)$	18
T_4	$(a, 10), (c, 6), (e, 6), (g, 5)$	37
T_5	$(b, 4), (c, 2) (e, 3), (g, 2)$	21

Example: $X = \{b, c\}$


its average utility $au(X) = (40 + 20 + 21) / 3 = 27$

its average cost $ac(X) = (10+1+8+3+4+2) / 3 = 9.3$

its support $s(X) = 3$

The LCIM algorithm

- An **algorithm** that starts from itemsets containing **single items** and recursively combine them to generate larger itemsets.
- Performs a **depth-first search**.
- **Applies search space pruning properties.**



Search space pruning using the support

This is a well-known properties in frequent itemset mining:

Property 1 (Support pruning). For any two itemsets $X \subseteq Y$, $s(X) \geq s(Y)$.

Search space pruning using the cost

Definition 3 (Lower bound on the cost). Let there be an itemset X , which appears in the transactions $g(X) = \{T'_1, T'_2, \dots, T'_N\}$. The **sequence of cost values of X** is the unique sequence $A(X) = (a_i)_{i=1}^N$ where $a_i = c(X, T'_i)$. Let $\text{sort}(A(X))$ be the unique non decreasing sequence $B = (b_i)_{i=1}^N$ satisfying $\forall i \leq K \leq N \quad |\{i \in \mathbb{N} | a_i = a_K\}| = |\{i \in \mathbb{N} | b_i = a_K\}|$. The **K largest cost values of X** is the sequence $A(X)^{(K)} = (c_i)_{i=1}^K$ of real numbers satisfying $\forall 1 \leq i \leq K, c_i = b_{N-K+i}$. The **average cost bound (ACB)** of X is defined as $acb(X) = \frac{\sum_{c_i \in A(X)^{(\minsup)}} c_i}{s(X)}$.

Property 2 (lower bound of the ACB on the average cost). For any itemset X , the average cost bound of X is a lower bound on the average cost of X . In other words, $acb(X) \geq ac(X)$.

Property 3 (anti-monotonicity of the ACB). For any two itemsets $X \subseteq Y$, then $acb(X) \leq acb(Y)$.

Property 4 (Search space pruning using the ACB). For an itemset X , if $acb(X) > \text{maxcost}$, X and its supersets are not low cost itemsets.

The cost-list data structure

- To facilitate the calculations, **each itemset** has a special structure called **cost-list**.
- **Example:**

The cost-lists of $\{a\}$, $\{b\}$ and $\{a, b\}$

$L(\{a\})$	
<i>utility</i>	95
<i>cost</i>	20
<i>tids</i>	$\{T_1, T_3, T_4\}$
<i>costs</i>	$\langle 5, 5, 10 \rangle$

$L(\{b\})$	
<i>utility</i>	81
<i>cost</i>	22
<i>tids</i>	$\{T_1, T_2, T_5\}$
<i>costs</i>	$\langle 10, 8, 4 \rangle$

$L(\{a, b\})$	
<i>utility</i>	40
<i>cost</i>	15
<i>tids</i>	$\{T_1\}$
<i>costs</i>	$\langle 15 \rangle$

Algorithm 1: LCIM

input : D : a transaction database,
 $\text{minsup}, \text{minutil}, \text{maxcost}$

output : all the low cost itemsets

- 1 Calculate the support of all items by reading the database;
- 2 Store the frequent items in a set ω ;
- 3 Establish the order of support ascending values on ω , and denote it as \succ ;
- 4 Create the cost-list of all frequent items in ω ;
- 5 Put each item i from ω in a set γ if $\text{acb}(\{i\}) \leq \text{maxcost}$ according to $L(\{i\})$;
- 6 Search $(\gamma, \text{minsup}, \text{minutil}, \text{maxcost})$;

Algorithm 2: Search

input : P : some itemsets with their cost-lists,
 $\text{minsup}, \text{minutil}, \text{maxcost}$

output : a set of low cost itemsets

- 1 **foreach** itemset $X \in P$ **do**
- 2 **if** $\text{au}(X) \geq \text{minutil} \wedge \text{ac}(X) \leq \text{maxcost}$ **then** Output X ;
- 3 Initialize X_{Extend} as the empty set;
- 4 **foreach** itemset $Y \in P$ that can be joined with X **do**
- 5 $Z \leftarrow X \cup Y$;
- 6 $L(Z) \leftarrow \text{Construct } (L(X), L(Y))$;
- 7 **if** $s(Z) \geq \text{minsup} \wedge \text{acb}(Z) \leq \text{maxcost}$ **then**
- 8 $| X_{\text{Extend}} \leftarrow X_{\text{Extend}} \cup \{Z\}$;
- 9 **end**
- 10 **end**
- 11 Search $(X_{\text{Extend}}, \text{minsup}, \text{minutil}, \text{maxcost})$;
- 12 **end**

Algorithm 3: Construct procedure

```
input :  $L(X), L(Y)$ : the cost utility list of two itemsets  $X$  and  $Y$ 
output : the cost-list  $L(Z)$  of  $Z = X \cup Y$ 


1 Initialize a cost-list  $L(Z)$  such that  $L(Z).utility = 0$ ,  $L(Z).tids = \emptyset$ , and  $L(Z).costs = \emptyset$ ;
2 foreach transaction  $T_w \in L(X).tids$  do
3   if  $\exists T_w \in L(Y).tids$  then
4      $L(Z).tids \leftarrow L(Z).tids \cup \{T_w\}$ ;
5      $L(Z).utility \leftarrow L(Z).utility + u(T_w)$ ;
6      $L(Z).costs \leftarrow \text{Merge}(L(X).costs, L(Y).costs)$ ;
7   end
8 end
9  $L(Z).cost = \sum L(Z).costs$ ;
10 return  $L(Z)$ ;
```

Two optimizations. We also include two performance optimization: (1) *Matrix Support Pruning* (MSP) consists of precalculating the support of all pairs of items in the initial database scan. Then, two itemsets X and Y are not joined as $X \cup Y$ if their last items have a joint support below $minsup$. (2) Efficient List Construction (ELC) the construct procedure is stopped as soon as there are not enough transactions left in $L(X).tids$ to attain $L(Z).tids \geq minsup$.


Experiments

- **Four datasets**
 - **Chess** has 3,196 transactions, 37 distinct items, and an average transaction length of 75 items.
 - **Mushroom** contains 8,416 transactions, 119 distinct items, and the average transaction length is 23 items.
 - **Accidents** contains 340,183 transactions with 468 distinct items, and an average transaction length of 34 items.
 - **E-Learning** contains 54 transactions, 11 distinct items, and the average transaction length is 3.8 items.
- **LCIM vs baseline** (LCIM with no optimizations)

Performance evaluation (1)



Performance evaluation (2)



Observations

- **LCIM** is not always faster than the **baseline**
- **Increasing maxcost** → execution time increases
pattern count increases
- **Increasing minsup** → execution time increases
pattern count increases
- **minutil** → negligible effect on execution time

Patterns found in e-learning data

- E-learning dataset describe activities done by 54 students during Session 4 of an e-learning environment, and has cost and utility values.
- We found several interesting patterns with **minutil = 10**, **maxcost = 100** and **minsup = 0.2%**

No	Patterns	Avg. Util	Avg. Cost	Sup	Trade-off
1	<i>Deeds_Es_4_3, Deeds_Es_4_2</i>	12.53	21.23	13	1.69
2	<i>Deeds_Es_4_4, Deeds_Es_4_1</i>	14.07	25.21	14	1.81
3	<i>Deeds_Es_4_4, Deeds_Es_4_5</i>	12.42	28.92	14	2.38
4	<i>Deeds_Es_4_1</i>	14.21	13.82	23	0.98
5	<i>Deeds_Es_4_1, Deeds_Es_4_2</i>	13.64	26.71	14	1.96
6	<i>Deeds_Es_4_1, Deeds_Es_4_2, Deeds_Es_4_5</i>	12.76	60.0	13	4.76
7	<i>Deeds_Es_4_1, Deeds_Es_4_5</i>	12.5	34.12	16	2.77
8	<i>Deeds_Es_4_2</i>	13.26	10.60	23	0.80
9	<i>Deeds_Es_4_2, Deeds_Es_4_5</i>	12.47	27.68	19	2.22
10	<i>Deeds_Es_4_5</i>	11.62	17.20	24	1.49

A pattern's **trade-off** is the ratio of its average cost to its average utility

Conclusion

Contributions

- A new problem of **low-cost high utility itemset mining**
- A new algorithm **LCIM**
- Search space pruning strategies and optimizations

Research opportunities

- Alternative algorithms and additional optimizations
- Extending the problem (top-k, closed, maximal itemsets...)



Open source Java data mining software, 240 algorithms
<http://www.phillippe-fournier-viger.com/spmf/>